Category	Check	Points	Total		
Required:					
Documentation on how to launch your work (10 points)	☑	10	10		
Choose your own adventure:					
Kafka Consumer runs in a container and stores data in (you can only choose one):					
a CSV (persisted outside the container) – 5 points	☐	5			
sqlite (persisted outside the container) – 10 points	☑	10			
postgres (in a separate container) --15 points	☐	15	10		
SparkML trained model – 10 points	☑	10	10		
Model retrains from data above – 10 points	☑	10	10		
Model is offered in a container – 10 points	☑	10	10		
Model is refreshed/redeployed by retraining above – 20 points	☑	20	20		
Model is refreshed/redeployed by retraining above w/o downtime – 10 points	☐	10	0		
A prediction is made via the model and a dataset sent to it – 5 points	☑	5	5		
A prediction is made via the model and windows of the data streamed from Kafka – 10 points	☑	10	10		
Prediction is correct > 95% of the time (based on our data not whatever you used as test data) – 15 points	☐	15	0		
Graphs:					
Web page serving graphs – 10 points	☑	10	10		
A graph visualizing the events being streamed – 10 points	☐	10	0		
A graph visualizing past outbreaks from the data we provided or you persisted (or both) – 10 points	☑	10	10		
A graph visualizing the next predicted outbreak event – 15 points	☐	15	0		
Other graphs – 5 points each	☑	5	5		
Everything you produce runs from a Makefile – 10 points	☑	10	10		
Everything you produce runs as a docker compose file – 10 points	☑	10	10		
Everything you produce has docker compose health checks – 10 points	☑	10	10		
Everything you produce has docker compose health checks and will restart on fail – 10 points	☑	10	10		
Everything you produce runs from a python or shell script – 5 points	☐	5	0		
The things you produce have tests to prove they work using pytest (these stack, i.e. 100% == 20 points)					
25% – 5 points	☑	5	5	We kind of have tests	
50% – 5 points	☑	5	5		
75% – 5 points	☐	5	0		
100% – 5 points	☐	5	0		

Everything you produce runs in Kubernetes (recommend kind) – 25 points	☑	25	25		
Documentation					
How to configure the model to be retrained more / less often – 10 points	☐	10	0		
	Total	Possible	Achieved		
		265	185		